

Problem 2

Minimizing the L2 penalized logistic regression cost function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) + \frac{\lambda}{2m} \sum_{j=1}^d \theta_j^2$$

1. (True/False) $J(\theta)$ has multiple locally optimal solutions.
False, because when minimized, $J(\theta)$ will have a unique global maximum because the Hessian is positive definite.

2. (True/False) Let $\theta^* = \operatorname{argmin}_{\theta} J(\theta)$ be a global optimum. θ^* is sparse.
False, θ^* is not sparse when $\theta^* = \operatorname{argmin}_{\theta} J(\theta)$ is a global optimum because the L2 penalized logistic regression just minimizes the values of theta so that they become close to 0, but don't actually become 0 unless the regularization term is infinity. This is due to the nature of the error term as a squared error (quadratic) will start taking smaller step sizes towards zero once the weights approach zero, therefore never actually reaching zero but can result in theta terms that are very small.

3. (True/False) If the training data is linearly separable, then some coefficients θ_j might become infinite if $\lambda = 0$.
True, if the training data is linearly separable, then the data should be able fully separated by an infinite number of lines (hyperplanes) graphically. The θ_j coefficients should be able to map to that line and the cost function, when minimized, should accurately point to all these hyperplanes. This means that the θ_j coefficients can map to an infinite number of hyperplanes, of which some coefficients θ_j could very well be infinite.

4. (True/False) The first term of $J(\theta^*)$ always increases as we increase λ .
True, if you ignore the case of $\lambda = 0$, then whenever λ is increased, the first term of $J(\theta)^*$ will always increase as well. As λ increases, the entire loss function will increase, continuing to increase with the addition term indefinitely. The first term of $J(\lambda)$ represents the data loss of the logistic regression model. As you increase λ , the data loss will increase as it is attempting to decrease the overfitting (as overfitting too much would actually cause the cost function in the training set to become 0), therefore λ increasing will cause

the first term of $J(\theta)$ to always increase.

kNN Classifier

Classifying test data with a knn-classifier

Question: Notice the structured patterns in the distance matrix, where some rows or columns are visible brighter. (Note that with the default color scheme black indicates low distances while white indicates high distances.)

What in the data is the cause behind the distinctly bright rows?

What causes the columns?

Answer: The data behind distinctly bright rows is a test sample that is high distance from every "category" within the training data. Therefore, it might be an object that is not currently categorized within the dataset or is very different from all other objects within the dataset.

The data behind distinctly bright columns are training set samples that do not seem to exist within the test set, and therefore show high distances from every item of the test set.

Choosing k by cross-validation

Chosen k value: 10

Obtained accuracy: 0.282